

# Manuscript Title

This manuscript ([permalink](#)) was automatically generated from [computer-aided-biotech/better-cb@faf4795](#) on February 15, 2021.

## Authors

---

- **Benjamín J. Sánchez**

 [0000-0001-6093-4110](#) ·  [BenjaSanchez](#) ·  [BenjaSanchez](#)

Department of Bioengineering, Technical University of Denmark, Kgs. Lyngby, 2800, Denmark

- **Daniela C. Soto**

 [0000-0002-6292-655X](#) ·  [dcsoto](#)

Genome Center, MIND Institute, and Department of Biochemistry & Molecular Medicine, Davis, CA 95616, USA

# Abstract

---

This should be the abstract.

## Introduction

---

Since Margaret Dayhoff pioneered the field of bioinformatics back in the sixties, the application of computational tools to the field of biology has vastly grown in scope and impact. Nowadays, biotechnological and biomedical research are routinely fed by the insights arising from novel computational approaches, machine learning algorithms and mathematical models. The ever increasing amount of biological data and the exponential growth in computing power will amplify this trend in the years to come.

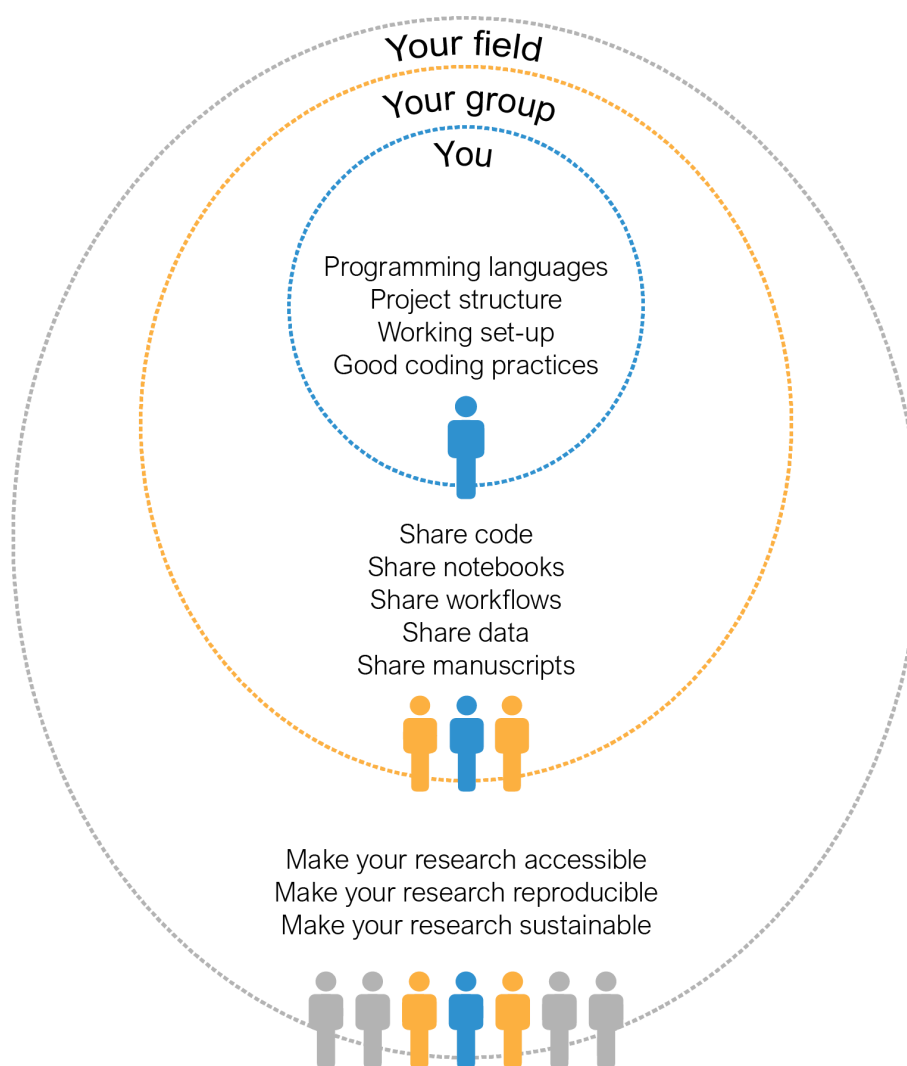
The use of computers to address biological matters encompasses a wide array of applications usually grouped under the terms of “computational biology” and “bioinformatics”. Although distinct definitions have been delineated for each one [\[1,2\]](#), here we will consider both under the umbrella term “computational biology”, alluding to any application that involves the intersection of computing and biological data. As such, a computational biologist can be a data analyst, a data engineer, a statistician, a mathematical modeler, a software developer, and many others. In praxis, the modern computational biologist will be a “scientist of many hats”, taking on several of the duties listed above. But first and foremost, we will consider a computational biologist as a scientist whose ultimate goal is answering a biological question or addressing a need in the life sciences, by means of computation.

Scientific computing requires following specific practices to enable shareable, reproducible and sustainable outputs that stand the test of time. Computing-heavy disciplines such as software engineering and data science have already adopted practices addressing the need for collaboration, visualization, project management, and strengthening of online communities. However, as a highly interdisciplinary and evolving field, computational biology has yet to acquire a set of universal “best practices”. As most computational biologists come from diverse backgrounds and oftentimes lack formal computational training, the absence of guidelines can lead to disparate and unsustainable practices that hinder reproducibility and collaboration, and slow down biomedical and biotechnological research.

Over the last decade, several researchers have published advice directed to bench scientists starting in either scientific computing [\[3,4\]](#) or computational biology [\[5\]](#). The advice encompasses a wide range of topics ranging from programming to project organization to manuscript writing. Other works have adopted a different approach, fixating in one powerful tool and diving deeper into its scope and applications, such as the software development and version control cloud service GitHub [\[6\]](#) and the web-application Jupyter Notebooks [\[7\]](#). Similarly, recent works have chosen to comprehensively address one specific need in computational biology such as workflow automation [\[8\]](#) and software library development [\[9\]](#). Although this advice proves immensely helpful, several aspects of the computational biology “journey” remain uncovered. Specifically, there is a lack of a clear roadmap regarding best practices from fundamental to advanced topics, in addition to practical examples tackling the complexities of this kind of research.

We premise that best practices in computational biology lie within a continuum that traverses three “levels”: the individual’s internal practices, the collaborative practices of a team, and the practices that allow a broader scientific community to access and engage over time with the research (Figure [1](#)). Each one of these levels has a different set of needs and challenges. Here, we compile a selected list of relevant practices and related tools advisable for each one of these levels, emphasizing their time

and place in a computational biology research project. Finally, we illustrate the utility of these practices in three case studies covering the broader spectrum of computational biology research.



**Figure 1:** The three “levels” of computational biology include your personal research, your group and collaborators, and your scientific field.

## Level 1: Personal Research

The computational biology “journey” begins with you—specifically, with the set of skills, tools and practices that you have put in place to conduct your research. Taking the time to optimally establish these building blocks will have high payoffs later, when you need to check your previous research (which will certainly happen). Consider that your most important collaborator is your future self, either of tomorrow or ten years from now. To consider all the aspects that make up a solid ground for any computational biology project, we devised a framework involving four main sequential steps (Table 1).

### Step 1: Choose your programming languages

Different programming languages serve distinctive purposes and have different idiosyncrasies. As such, choosing the right programming language for a project depends on your research goals and personal preference or skill. Additionally, communities usually favor the usage and training of some

programming languages over others; utilizing such languages may facilitate integrating your work within the existing ecosystem.

As computational biology becomes a data intensive discipline, interacting with high-performance computing (HPC) operating systems has become a hallmark of the field. HPC infrastructures commonly use Unix/Linux distributions as their operating system. To interact with these platforms, you need to use a the command line interpreter known as the “Unix shell”. There are multiple versions of Unix shells, being Bash one of the most widely adopted. Besides providing an “user interface”, the shell is also a scripting language that allows manipulating files and executing programs through “shell scripts”. Unix/Linux operating systems have other interesting perks, such as powerful and fast commands for searching words and manipulating files (e.g. `sed`, `grep` or `join`) as well as the language AWK, that can perform quick text processing, including arithmetic operations.

One of the most common task of any computational biologist is data analysis. The job of a data analyst involves data cleaning, exploration, manipulation, and visualization. Currently, Python is the most widely used programming language for data analysis worldwide [10,11]. Computational biology research has followed this trend, making Python one of the most popular languages among researchers. As machine learning and deep learning are more widely adopted in biological research, Python usage will keep growing. Python usage has been facilitated by the availability of packages for biological data analysis accessible through package managers such as Pip or Conda. Similarl to Python, R is the other prominent language in the field. Arguably, one of R main strengths is its wide array of tools for statistical analysis. Of particular interest is the Bioconductor repository, where many gold-standard tools for biological data analysis have been published. R usage in data science has deeply benefited from the Tidyverse packages, increasing the readability of the R syntax for both data manipulation via `dplyr`, and visualization via `ggplot2`.

Oftentimes, computational biologists require coding their own sets of instructions for processing data via scripts or programs. A script can be described as lightweight single-file program developed to tackle a narrow purpose. They are likely written in an interpreted programming language instead of a compiled one. Interpreted programming languages execute the program directly, without previous compilation, meaning each statement is run individually. They are quick to edit and can be run interactively, at expense of computational performance. In computational biology, the current most common multi-purpose scripting languages are Python and R. When working in a HPC, Shell/Bash scripting is also widely used. Perl and Matlab are also a popular language among bioinformatics and systems biology, respectively. A program, in the other hand, is a larger tool that usually combines multiple scripts and works as a “black box” to the user. It is designed to tackle computationally intensive problems, thus, a compiled language is preferred. Several tools designed for high-weight biological data processing have been written in C/C++. In recent years, however, scientists have been turning to Rust because of its speed, safety and friendly community [12]. If computational performance is not a concern, Python and R are suitable alternatives for coding programs.

Biological data processing is rarely a one-step process. To go from raw data to useful insights, several steps need to be taken in a specific order, accompanied by a plethora of decisions regarding parameters. Computational biologists have addressed this need by embracing workflow management systems to automate data analysis pipelines. A pipeline can be written as a Shell script where a handful of commands are written one after another, using Shell variables and Shell scripting syntax when needed. Although effective, this approach provides little control over the workflow, and lacks features to run isolated parts of the pipeline or track changes in input and output files. To overcome these limitations, a Bash script can be “upgraded” using the GNU Make program, which was originally designed to automate compilation and installation of software, but that it is flexible enough for building other types of workflows. Nowadays, however, several dedicated bioinformatics workflow managers have been developed. Snakemake is a workflow management system written in Python, allowing to incorporate the syntax of the tool with standard Python code. Similarly, Nextflow was build

as an extension of the Groovy—a programming language for the Java virtual machine—and can execute Groovy code. These tools not only provide control over any step of the pipeline, but also offer features like interacting with job schedulers, software environment managers, and cloud computing support. Alternatively, workflows can be written in the Common Workflow Language (CWL)—a declarative standard to define workflows with the goal of enabling portability and reproducibility. Workflows written in CWL can be run in any CWL-enabled engine.

## Step 2: Define your project structure

After choosing your programming languages and before starting any coding, we advise you to come up with a well-thought-out project structure. This design should be intentional and tailored to the present and future needs of your project—remember to be kind to your future self. A computational biology project requires, at the very least, a folder structure that supports code, data, and documentation. Although tempting, cramming all kind of files in a unique folder is unsustainable. Instead, separate each one on different folders and subfolders if needed. As additional principles consider documentation as a requirement and raw data as immutable. To simplify this process, you can base your project structure on research templates available off-the-rack. For data science projects, the python package Cookiecutter Data Science cut down the effort to the very minimum [13]. Running the package prompts a questionnaire in the terminal where you can input the project name, authors and other basic information. Then, the program generates a folder structure to store data—raw and processed—separated from notebooks and source code, as well as pre-made files for documentation such as a “readme”, a docs folder and a license. Similarly, the Reproducible Research Project Initialization (rr-init) offers a template folder structure that can be cloned from a GitHub repository and modified by the user [14]. Although the later is simpler than the former, both follow an akin philosophy aimed at research correctness and reproducibility [15]. For workflow automation projects, we advise a similar folder structure. Snakemake recommends to store each workflow in a dedicated folder separated into workflow-related files—the Snakefile, rules and scripts—results and configuration [16]. In all cases, the folder must be initialize as a git repo for version control (See Step 4).

Beyond files and folders, the software and dependencies needed to run an analysis, workflow or program, are also part of the project structure itself. The intricacies of software installation and dependency management are not to be underestimated. Fortunately, package and virtual environment managers significantly reduce this burden. A package manager is a system that automates the installation, upgrading, configuration and removing of community-developed programs; a virtual environment manager, in the other hand, is a tool that generates isolated “environments” containing programs and dependencies that are functionally independent from other environments or the default operating system. Once a virtual environment is activated, a package manager can be used to install third-party programs.

We believe that a computational biology project must start with its own virtual environment. The main reason is reproducibility: environments save the project's dependencies and can restore them at will so the code can be run in any other computer. There are multiple options for both package and virtual environment management—some are language-specific; others, language-agnostic. If you are working with Python, you can initialize a Python environment using `virtualenv` or `pipenv` (where different Python versions can be installed). Inside the environment, you can use the Python package manager `Pip` to add Python code submitted to the Python Package Index (PyPI), GitHub or locally. For the R language, R-specific environments can be created using `renv`. After initializing the environment, packages can be installed via `install.packages` function from the Comprehensive R Archive Network (CRAN) and CRAN-like repositories. R also has `BiocManager` to install packages from the Bioconductor repository, which contains relevant software for high-throughput genomic sequencing analysis. To fully manage a dependencies beyond installation, RStudio developed the RStudio Package Manager which works with third-party code available in CRAN, Bioconductor, GitHub

or locally. A language-agnostic alternative is Conda—an increasingly popular package manager and a virtual environment manager. It supports program installation from the Anaconda repository, which contains the channel Bioconda specifically tailored to bioinformatics software. Also, if Python is installed, Python dependencies can be installed via pip. Conda is particularly helpful when working with third-part code in all sorts of languages—a common predicament for the computational biologist. Conda package and environment manager is included in both Miniconda and Anaconda distributions. The former is a minimal version of Anaconda, containing only Conda, Python, and a few useful packages.

### **Step 3: Choose your working set-up**

With the foundation in place, the next step is to start coding. However, a more practical question needs to be answer first: where to code. The simplest tool available for this purpose are text editors. Since writing code is ultimately writing text, any tool where characters can be typed fulfills this purpose. However, coding can be streamlined by additional features as those found in code editors—text editors especially developed for writing code. Crucial features to facilitate coding include syntax highlight, indentation or autocompletion. Commonly used desktop editors include Atom, Sublime, Visual Studio Code, and Notepad++ (Windows only), all which a myriad of plugins available to enhance the coding experience. Command line text editors are also suitable options for coding, being Vim and Emacs the most powerful ones. All of these tools share the advantage of being language agnostic, allowing easy switching between languages, especially handy for the polyglot computational biologist.

In addition to text editors, integrated development environment (IDE) are also popular options for coding. As its essence, IDE are supercharged text editors, i.e. with multiple other features that make writing code easier. The main parts of an IDE are a code editor (with syntax highlight, indentation and suggestions), a debugger, a folder structure, and a way to execute your code (a compiler or interpreter). IDEs are not language agnostic, meaning that they allow to code in one language. The array of features also comes at a cost—IDEs usually use more memory and imply more visual clutter, if that is a concern of yours. For Python, Jupyter Lab, Spyder and JetBrains' PyCharm are popular options, while for R, RStudio is the gold-standard. Notably, the differences between an IDE and a code editor are somewhat blurry, especially when enough plugins have been added to a code editor.

In the latest years, notebooks have acquired relevance in computational biology research. A notebook is an interactive application that combines live code (read-print-eval loop or REPL), narrative, equations and visualizations. Common notebooks use an interpreted languages such as Python or R, and narrative follows markdown syntax. Data analysis greatly benefits from using notebooks instead of plain text editors or even IDEs: the combination of visuals and texts allows researcher to tell compelling stories about their data, and the interactivity of its code enables quick testing of different strategies. Jupyter notebook is a popular web-based interactive notebook developed originally for Python coding, but also accepts R and other programming languages upon installation of their kernels—a computing engine that executes the notebook's live code “under the hood”. Jupyter notebook can also be run in the cloud using platforms such as Google CoLab and Amazon WebServices, taking advantage of the current trend of cloud computing. RStudio also allows the generation of R-based notebooks known as R Markdown, which is specially well-suited for generating reports.

### **Step 4: Follow good coding practices**

After dealing with steps one to three, finally comes writing code. Coding, however, requires good practices to ensure correctness, sustainability and reproducibility for you, your future self, your collaborators (as we will discuss in Level 2) and the whole community (as we will discuss in Level 3). First and foremost, you need to make sure your code works correctly. In computational biology, correctness implies biological and statistical soundness. Both are big topics beyond the scope of this



manuscript. To achieve the former, however, a useful approach is to design positive and negative controls in your program, analysis or workflow. In an experiment, a positive control is a control group that is expected to produce results; a negative control is expected to produce no results. The same approach can be applied to computation, using input data whose output is previously known. Biological soundness can also be tested by quickly assessing expected orders of magnitude in both, intermediate and final files.

Beyond the correct functioning of your code, you will need to pay attention to the way your code looks, also known as “coding style”. This includes a series of small and ubiquitous decisions regarding where and how to add comments; indentation and white spaces usage; variable, function and class naming; and overall code organization. It is true that, as in writing, there is a lot of your own personality in the way you code. However, sticking to existing coding style rules facilitates collaborations with your future self and others. Indeed, as we sometimes have trouble reading our own handwriting, we can also struggle reading our own code if we overlook any guidelines. At the very least, your code must display internal consistency. Even better, you can follow any of the multiple coding style rules that have been published. Although arbitrary, most of these rules have been developed with readability in mind. A good place to start, however, are style guides from software development teams. Google, for example, has published guidelines for Python, R, Shell, C++, and HTML/CSS [17]. Also, a series of guidelines for the Python programming language have been published with the name of Python Enhancement Proposal (PEP), where the most widely adopted is PEP-8 [18]. To aid flagging stylistic errors in your code, tools called “linters” are usually included with code editors and IDEs or provided as plugins.

In the matter of code styling, two topics merit additional attention: variable naming and comments. Variable names should be descriptive enough to convey an idea about the variable, function or class’ content and use. The goal is to produce “self-documented” code that reads close to plain English. To do so, use multi-words variable names if necessary. In such cases, the most common conventions include Camel Case, where the second and subsequent words are capitalized (“camelCase”); Pascal Case, where all words are capitalized (“PascalCase”); and Snake Case, where words are separated by underscores (“snake\_case”). Notably, all these conventions can be used in a same coding style to differentiate variables, functions and classes. For example, PEP-8 recommends Snake Case for functions and variables, and Pascal Case for class names. In addition to master variable naming, code comments—explanatory human-readable statements not evaluated by the program—are necessary to enhance the code’s readability. No matter how beautiful and well-organized your code is, high-level code decisions will not be obvious unless stated. As a corollary, code explanations that can be deduced from the syntax itself should be omitted. Comments can span a single line or several ones, forming a block, and can be found in three strategic parts: at the top of the program file (“header comment”), which contains the author and the date of the code and what it accomplishes; above every function (“function header”), which contains the purpose and behavior of the function; and in line, next to tricky code whose behavior is not obvious or warrant a remark.

When working with a sizable code base, a good practice related is to strive for modularity—splitting your code’s functionalities into independent entities known as modules. Modularity enhances code readability and reusability—enabling your code to be used by you or others in future applications—and expedites maintenance. In Python, subdivisions are defined as follow: a module is a collection of functions and global variables, a package a collection of modules, a library a collection of packages, and a framework a collection of libraries. Modules are simply files with the `.py` extension. Packages, in the other hand, must be indicated to the Python interpreter adding a file named `__init__.py` (which could be empty or not).

Code styling rules also apply to data science notebooks. However, when writing notebooks you must also engage in “literate programming”—a programming paradigm where the code is accompanied by human-readable explanation of its logic and purpose. In other words, notebooks must tell a story

about the analysis, connecting the dots between the code, the results and the figures. Human-readable language is often written in Markdown—a lightweight markup language. Little has been written about good practices for literate programming, but we advise you to explain the purpose of each chunk of code and provide some interpretation of its results.

Equally as important as writing good code is to use version control—the practice of tracking and managing changes in your code. This is a good and necessary practice even if your only collaborator is your future self. One of the main advantages of version control is keeping a change log of your files that can be utilized to go back to previous versions of your code, and remind you of previous approaches disregarded in newer versions. The most widely used version control system, Git, achieves these tasks with the command `git checkout`. Additionally, version control also allows you to safely try new functionalities using “branches”—carbon copies of the main original branch (known as “master”) where you can add code independently and optionally merge it to the original one. Git creates branches using the command `git branch` and the same `git checkout` can be used to switch among them. We will discuss the utility and implementation of branches in collaborative projects in the next section (See Level 2: Collaboration). Nowadays, there are multiple code repositories that also provide version control with Git, such as GitHub, GitLab or Bitbucket. They have the additional benefit of backing up your code and code history in the cloud, keeping your work safe and shareable. You can interact with these platforms using the browser or via graphic user interfaces (GUI) such as GitHub Desktop or GitKraken.

**Table 1:** Steps to start any computational biology project.

| Step   | Use case                          | Common tools   |
|--|-----------------------------------|--|
| <b>Step 1:</b> Choose your programming languages | Interacting with a Unix/Linux HPC | • Shell/Bash   |
|  | Data analysis                     | • Python, R  |
|  | Scripts and programs:             | • Interpreted: Python, R, Perl<br>• Compiled: C/C++, Rust  |
|  | Workflows                         | • Snakemake (Python), Nextflow (Groovy), CWL   |
| <b>Step 2:</b> Define your project structure     | Project structure                 | • Templates: Cookiecutter Data Science, rr-init<br>• Workflows: Snakemake structure                            |
|  | Virtual environment managers      | • Language-specific: pipenv (Python), virtualenv (Python), renv (R)<br>• Language agnostic: Conda              |
|  | Package managers                  | • Language-specific: pip (Python), BiocManager (R), R Studio package manager (R)<br>• Language-agnostic: Conda |
| <b>Step 3:</b> Choosing your working set-up      | Text editors                      | • Desktop applications: Atom, Sublime, Visual Studio Code, Notepad++<br>• Command line: Vim, Emacs             |
|  | IDEs                              | • For Python: Jupyter Lab, JetBrains/PyCharm, Spyder<br>• For R: R Studio                                      |
|  | Notebooks                         | • Jupyter (Python, R), R Markdown (R)  |
| <b>Step 4:</b> Follow good coding practices      | Coding style                      | • Styling guides: PEP-8 (Python), Google (Python, R)<br>• Linters  |
|  | Literate programming              | • Markdown   |
|  |                                   |  |



| Step | Use case        | Common tools   |
|------|-----------------|--|
|      | Version control | <ul style="list-style-type: none"> <li>• Version control system: Git</li> <li>• Code repositories: GitHub, GitLab, Bitbucket</li> <li>• Git GUIs: GitHub Desktop, GitKraken</li> </ul> |

## Level 2: Collaboration

Collaboration is a key aspect of scientific research, but it is especially relevant in computational biology, where interdisciplinary knowledge is often needed. Collaborators can take different forms: your boss or advisor, colleagues or lab mates, other laboratories, people from academia or industry, or your future self (as discussed in Level 1). Although collaborators can have a wide range of involvement with your project—from co-authors to commenters—they all share a direct relationship with you and your research, comprising a group of a dozen of people at most (contrary to a community, which is an open group of a large number of people, as we will discuss in Level 3). Each type of collaboration requires its own set of good practices, which we will cover in the next paragraphs.

### 2.1 Share code

Sharing code is one of the most common practices in software development, where large teams work together developing highly complex functions and scripts. Although computational biology projects are usually not as big, proper ways of sharing code are still essential, as it is not desired to have file conflicts as soon as two different researchers change the same piece of code. As mentioned in the previous section, hosting services such as GitHub [19], GitLab [20] and Bitbucket [21] (Table ??) allow for having a Git repository stored online, by creating a copy of the repository known as the “remote”, which becomes the official version of the repository. The key advantage of using a remote is that there will be no direct interaction between different local copies of the repository, also known as a “clones”, but instead each clone only will interact with the remote, and can only update the remote if there are no conflicts. This way, if a collaborator updated the repository with some changes, another collaborator will not be able to send their changes until they make sure to update their local copy with the changes already present online.

Tools for collaborative research. {#tbl:collaboration-tools}

| Goal                         | Tools   |
|------------------------------|---|
| Share code                   | <ul style="list-style-type: none"> <li>• <i>Hosting services</i>: <b>GitHub</b> [19], <b>GitLab</b> [20], <b>Bitbucket</b> [21].</li> <li>• <i>Git branching strategies</i>: <b>Github flow</b> [22].</li> <li>• <i>Tests</i>: correctness (e.g. <b>pytest</b> [23]), coverage (e.g. <b>codecov</b> [24]), automation (e.g. <b>tox</b> [25], <b>Travis CI</b> [26], <b>Github Actions</b> [27]).</li> <li>• <i>Code reviews</i>: <b>Github</b> [28], <b>Crucible</b> [29], <b>Upsource</b> [30].</li> </ul> |
| Share data                   | <ul style="list-style-type: none"> <li>• <i>FAIR principles</i> [31].</li> <li>• <i>Tidy data</i> [32].</li> <li>• <i>Data version control</i> [33].</li> </ul>   |
| Share data science notebooks | <ul style="list-style-type: none"> <li>• <i>Static</i>: <b>GitHub</b>, <b>GitLab</b>, <b>NBviewer</b> [34].</li> <li>• <i>Interactive</i>: <b>Binder</b> [35], <b>Google CoLab</b> [36].</li> <li>• <i>Comparative</i>: <b>ReviewNB</b> [37].</li> </ul>  |
| Share workflows              | <ul style="list-style-type: none"> <li>• <i>General hosting services</i>: <b>GitHub</b>, <b>GitLab</b>, <b>Bitbucket</b></li> <li>• <i>Dedicated workflow repositories</i>: <b>WorkflowHub</b> [38]</li> </ul>  |

| Goal              | Tools   |
|-------------------|---|
| Share manuscripts | <ul style="list-style-type: none"> <li>• <i>General-purpose word processors</i>: <b>Google Docs</b> [39], <b>Office 365</b> [40].</li> <li>• <i>Scholarly word processors</i>: <b>Authorea</b> [41].</li> <li>• <i>Online applications supporting Markup Languages</i>: <b>Overleaf</b> (LaTeX) [42], <b>Manubot</b> (Markdown + GitHub) [43].</li> </ul> |

In order to guarantee that different collaborators can work simultaneously in the same repository, a good idea is to implement some type of branching strategy in the repository (Table ??). In a small team of collaborators, the most common strategy is to have a single `master` branch and generate from it branches that each different developer can work on. Then, whenever the developer is ready, they can request to combine, or “merge”, the changes from their branch into the master branch, in a process known as “pull request”, or PR for short. Once a PR has been opened, collaborators can review it, and if it fulfills their criteria, approve it so that it can be merged into the master branch: any succeeding branches will now have those commits already included as part of their history. This branching strategy is sometimes referred to as Github flow [22] and will suffice for most projects. For more complex branching systems, see Section [3.3: Make your research sustainable](#).

Using Git hosting services for collaboration has many additional benefits. The commit history not only shows what was done at each point in time, but also which collaborator did it, so that if e.g. a bug was introduced, commands such as `git blame` will show which collaborator caused it. Collaborators can also create “forks”, i.e. full copies of repositories under their own possession, for e.g. having different a version of the software that works for a different purpose. Git hosting services can be accessed interactively online, or from the terminal with tools such as GitHub CLI [44]. Finally, Git hosting services also allow collaborators to open issues [45] for listing pending to-do’s and/or asking questions, acting as an open forum for development discussions, which has the advantage of remaining accessible for the future, if new collaborators would join the research project later (as opposed to closed e-mail discussions). We will discuss additional advantages of using Git hosting services, in terms of interacting with a user base, in [Level 3: Community](#).

Another important concept to internalize when developing code, especially together with other collaborators, is to develop unit tests (Table ??). Unit tests are scripts that will run to determine if specific modules/functions work as intended within the codebase, so that if later the function grows in scope, its proper basic functioning is ensured. For instance, if a function was defined for adding numbers, a simple test would be to asses if the function outputs 13 when the inputs 6 and 7 are provided. Tools such as `pytest` [23] for Python and `testthat` [46] for R exist to then detect said scripts, and run all of them to display if any specific section is failing. It is a very good practice to develop tests at the same time you develop code (at the personal research level), as adding tests *a posteriori* is significantly harder (but sometimes inevitable). Going beyond testing correctness, tools such as `flake8` [47] will test styling preferences (for complying with PEP8), `safety` [48] will test for vulnerabilities among the software’s dependencies, and `Codecov` [24] will test which percentage of the codebase is tested, given that as a rule of thumb, the more code lines tested, the more reliable a software is. All these different types of tests can be funneled into a single testing pipeline that can run automatically whenever desired. This process is known as Continuous Integration (CI), and can be tuned to run locally whenever commits are made, or online whenever a pull request is opened and/or merged. When running locally, an environment manager / command line tool such as `tox` [25] helps to ensure all tests are ran under e.g. different python versions. For setting up the CI cycle online, different dedicated CI tools such as Travis CI [26] or Circle CI [49] exist, and more recently GitHub actions [27] has been introduced for running the integration directly from Github.

Having tests is a great way of ensuring code fulfills a certain level of correctness and styling. However, it is no replacement for a human assessment to see if the code is correct, necessary and useful. Therefore, code reviewing is essential whenever developing code in collaboration (Table ??). Tools such as Crucible [29] and Upsource [30] exist for making in-line reviews of each file, but the most

common approach, if the software is stored in a repository, is to directly review using the online review tools from the hosting service. In the case of Github [28], this not only allows the reviewer to open a comment in any line of the code (which creates a thread for the original author to reply in), but also to suggest changes, for the author to in turn approve/dismiss. When reviewing, there are a series of things to look for (from functionality to documentation), and good practices to keep in mind (such as phrasing the comments in a constructive way), which are outside of the scope of this review but presented in detail elsewhere [50,51].

## 2.2 Share data

The practices of sharing data stem from the same place as with sharing code: we should store our dataset and any changes to it in a repository, and ensure it complies with standards by testing its quality. However, due to data having a more consistent structure than code, as data is often processed and outputted by machines in standard formats, there are additional criteria that should be considered when we share it with collaborators (and later on with the community). The main set of guidelines that represent these criteria were outlined a few years ago in what is known as the FAIR principles [31]: data should be Findable, i.e. easy to find online; Accessible, i.e. easy to access once found; Interoperable, i.e. easy to integrate with other data/applications/workflows/etc.; and Reusable, i.e. presented in a way that allows for others to use it for the same purpose or different settings.

For making data findable, research repositories such as Zenodo [52] and Figshare [53] allow you to assign a digital object identifier (DOI) to any group of files you upload, including data and/or code. Alternatively, regular code repositories like Github can be used instead, as you can use specific commits and/or releases to identify specific versions of the data (see [3.1. Make your research accessible](#)), in combination with extensions for Large File Storage such as git LFS [54], in the case of data files larger than 100 MB [55]. A final alternative is the Data Version Control (DVC) initiative [33], which is especially useful when doing machine learning, as it can keep track of data, machine learning models and even scoring metrics.

For making data accessible, we encourage as much as possible to make your repositories open access so that everything is accessible to everyone, but in cases in which you or your collaborators prefer some restrictions, you can create guest accounts to provide access to private repositories. For making data interoperable, distinctions between raw and clean data have been made [???], with the raw data being the same files that came out of the measuring device, and the clean data the files that are ready to be used for any computational analysis. An important characteristic that clean data should have is to be “tidy”, which is reviewed in detail elsewhere [32]. Finally, for making data reusable, thorough documentation of the data - including experimental design, measurements units and sources of error - is required.

## 2.3 Share data science notebooks

As we previously discussed, Jupyter Notebook have become a fundamental tool of data analytics. Accordingly, you will likely need to share your notebook with collaborators at some point. To do this, there are static and interactive options. The former, as the name indicates, share computational notebooks as a rendered text, written internally in HTML. Static notebooks are a good option when you want to avoid any modifications and can work as an archive of past analyses, but interacting with its content is cumbersome—the file must be downloaded and run in a local Jupyter installation. Git-based code repositories, such as GitHub and GitLab, automatically render notebooks that can be later shared pointing collaborators to the GitHub repository. To ease this process, the Project Jupyter provides a web application called NBviewer, where you can paste a Jupyter Notebook’s URL, publicly hosted in GitHub or elsewhere, and renders the file into a static HTML web page with a stable link.

Interactive notebooks, in the other hand, not only render the file but also allow collaborators to fully interact with it, tinkering parameters or trying new input data—no installation required. The Binder Projects (which is also part of the Project Jupyter) offers the Binder service, where any publicly hosted Git-based repository can be open with a Jupyter Notebook interface. The user can fully interact with any notebook within the repository, although changes will not be saved to the original file. The platform supports Python and R among other languages, and any additional packages required to run the analysis need to be specified in a configuration file within the repository. Similarly, Jupyter Notebooks can be run interactively using Google CoLaboratory (CoLab), which is available to anyone with a Google account. Notebooks can be updated locally, from any public GitHub repository, or from Google Drive, where also imported files are also saved. In both cases, the machines provided by these services are comparable to a modern laptop. Thus, these tools are not suitable for some computing-intensive tasks common to computational biology problems.

Besides sharing notebooks, oftentimes computational biologists need to work and edit a notebook together. In those cases, notebooks need to be treated as any other piece of code: updates from different collaborators must be managed with version control in a platform such as GitHub. The problem, however, is that Git-based hosting services deal with notebooks as if they were HTML text, where changes between versions are hard to visualize. To better compare these changes, there is NBreview, which renders and display in parallel the old and new versions of a notebook for easy comparison. The tool can be easily installed using your GitHub account, and notebooks can be reviewed from their website.

## 2.4 Share computational workflows

Computational biology projects often demands sharing multi-step analyses with dozens of third-party software and dependencies. Although these steps can be shared as documentation, complex workflows are better shared as stand-alone code that can be easily run with minimal file manipulation from collaborators. Doing so can safeguard the reproducibility and replicability of the analysis, leading to better science and less issues down the road.

The simplest way to share a pipeline is to generate a Bash script that receives input files from the command line, thus, allowing to run it with different input data. However, Bash scripts offer little control over the overall workflow and cannot re-run specific parts of the pipeline. To address these issues, pipelines are better shared using a workflow automation system. Theoretically, all the instructions regarding the workflow could be written in the main pipeline file—in Snakemake, the `.smk` file (or Snakefile); in Nextflow, the `.nf` file; and in CWL, the `.cwl` file. However, to ensure reproducibility, it is a good practice to share complete pipelines, meaning folder structures, additional files and software specification, as well all custom scripts developed for the analysis. These files can be shared using the same tools as other forms of code, namely GitHub or any other Git hosting services. Alternatively, they can be uploaded to hosting services specialized in workflows, like WorkflowHub [38], currently in beta.

When sharing workflows, consider that sharing software versioning is necessary for your collaborators to reproduce your pipeline using their own computing setup. Conda environments, for example, can be easily created from an environment file (in YAML language), which can be exported from an existing environment. Notably, Snakemake and Nextflow can be configured to automatically build isolated environments for each rule or step, enabling running different versions of a program within the same pipeline (which is especially helpful when needing both Python 2 and 3). Besides sharing the specifications of an environment, it is possible to share the environment itself via containers, using platforms like Docker and Singularity, which are especially helpful to share environments with a broader community, as we will discuss in Level 3.

## 2.5 Write manuscripts collaboratively

Writing articles is arguably the main way where we can share our research with the scientific community and for that purpose, the world. However, in a highly interdisciplinary field as computational biology, writing manuscripts is also a collaborative effort, where multiple people is directly involved in the crafting of a manuscripts. The traditional computer tools for writing documents, therefore, are oftentimes not suitable for this type of collaboration, resulting in files with different names jumping from one e-mail inbox to another, resulting in multiple and likely contradictory final versions. Let's avoid this by streamline collaborative manuscript writing with tools made for that purpose.

Big companies have become aware of the need for collaborative writing, developing online applications that can be simultaneously edited by multiple people. Google's GDocs and Microsoft's Office 365 are well-known word processors designed for this purpose, where the text is displayed with the exact appearance than in a printout (know as *What-You-See-Is-What-You-Get*, or WYSIWYG) and the text can be formatted making use of the internal features of the application. The advantage of these technologies is that they are extremely user-friendly, and require no additional knowledge. They are a good option when one or more of your collaborators seeks simplicity, but they are not specifically tailored for the needs of scientific writing, such as adding references, equations and figures. Fortunately, third-party companies have developed plugins for these applications to add references to your document. Companies like Authorea have developed their own online application specifically designed for writing manuscripts. Authorea, in particular, offers templates for different type of research projects, allows you to manage collaborators from the same platform, and easily add reference from using identifiers (DOI, PubMed, etc.). Consider, nonetheless, that some collaborators may not want to adopt a new tool exclusively for writing manuscripts.

In addition to word editors, text editors are a competitive option to write manuscripts when combined with a markup language—a human-readable computer language that uses tags to delineate formatting elements in a document that will be later rendered. Since the formatting process is internally handled by the application, styling elements (headers, text formatting, equations) can be easily written in text, even achieving greater consistency than word processors. Disciplines closely related to computational biology, such as statistics and mathematics, have historically used the markup language LaTeX for writing articles. This language has a specific syntax to write mathematics constructs as simple text, making it a sound choice for papers with lots of equations. To aid collaborative writing, platforms like Overleaf provide online LaTeX editors, supporting features like real-time editing. In addition to LaTeX, an emerging trend in collaborative writing is to use the lightweight markup language Markdown within the GitHub infrastructure. The software Manubot provides a set of functionalities to write scholarly articles within a GitHub repository, leveraging all the advantages of Git version control and the GitHub hosting platform [56]. For example, it provides cloud storage, version control, and facilitates the maintainers' work by managing updates via pull requests. The GitHub user interface also allows off-line discussions about the manuscript using issues" and task assignment (See level 3 for tips on project management). Manubot, in particular, accepts citations using manuscripts identifiers and renders automatically renders the article in PDF and HTML formats. As a drawback, it requires technical expertise in Git and familiarity with GitHub; as an upside, its reliable infrastructure scales well to large and open collaborative projects.

### Level 3: Community

---

Goal: How to develop and maintain a computational biology project with community feedback over time?

Topics: - GitHub releases and semantic versioning - Git Flow, GitHub Issues - Continuous integration and unit tests - Dependencies per user: pip-tools - Sharing software as Python packages, Conda/Bioconda or containers (Docker, Singularity) - Include license (MIT) and DOIs - Documentation: read the docs.

## Case studies

---

- Example of computational biology project 1: RNA-seq analysis (workflow)
- Example of computational biology project 2: Genome-scale metabolic model (systems biology project)
- Example of computational biology project 3: Software development (computational tool)

## Conclusion

---

Pending

## Acknowledgments

---

Pending



# References

---

## 1. NIH working definition of bioinformatics and computational biology

Huerta Michael, Downing Gregory, Haseltine Florence, Seto Belinda, Liu Yuan  
(2000-07-17)

<https://www.kennedykrieger.org/sites/default/files/library/documents/research/center-labs-cores/bioinformatics/bioinformatics-def.pdf>

## 2. What is bioinformatics? A proposed definition and overview of the field.

NM Luscombe, D Greenbaum, M Gerstein

*Methods of information in medicine* (2001) <https://www.ncbi.nlm.nih.gov/pubmed/11552348>

PMID: [11552348](https://pubmed.ncbi.nlm.nih.gov/11552348/)

## 3. Good enough practices in scientific computing

Greg Wilson, Jennifer Bryan, Karen Cranston, Justin Kitzes, Lex Nederbragt, Tracy K. Teal

*PLOS Computational Biology* (2017-06-22) <https://doi.org/gbkbwp>

DOI: [10.1371/journal.pcbi.1005510](https://doi.org/10.1371/journal.pcbi.1005510) · PMID: [28640806](https://pubmed.ncbi.nlm.nih.gov/28640806/) · PMCID: [PMC5480810](https://pubmed.ncbi.nlm.nih.gov/PMC5480810/)

## 4. Software engineering for scientific big data analysis

Björn A Grüning, Samuel Lampa, Marc Vaudel, Daniel Blankenberg

*GigaScience* (2019-05) <https://doi.org/gf4f4m>

DOI: [10.1093/gigascience/giz054](https://doi.org/10.1093/gigascience/giz054) · PMID: [31121028](https://pubmed.ncbi.nlm.nih.gov/31121028/) · PMCID: [PMC6532757](https://pubmed.ncbi.nlm.nih.gov/PMC6532757/)

## 5. So you want to be a computational biologist?

Nick Loman, Mick Watson

*Nature Biotechnology* (2013-11-01) <https://doi.org/p3j>

DOI: [10.1038/nbt.2740](https://doi.org/10.1038/nbt.2740) · PMID: [24213777](https://pubmed.ncbi.nlm.nih.gov/24213777/)

## 6. Ten Simple Rules for Taking Advantage of Git and GitHub

Yasset Perez-Riverol, Laurent Gatto, Rui Wang, Timo Sachsenberg, Julian Uszkoreit, Felipe da Veiga Leprevost, Christian Fufezan, Tobias Ternent, Stephen J. Eglén, Daniel S. Katz, ... Juan Antonio Vizcaíno

*PLOS Computational Biology* (2016-07-14) <https://doi.org/gbrb39>

DOI: [10.1371/journal.pcbi.1004947](https://doi.org/10.1371/journal.pcbi.1004947) · PMID: [27415786](https://pubmed.ncbi.nlm.nih.gov/27415786/) · PMCID: [PMC4945047](https://pubmed.ncbi.nlm.nih.gov/PMC4945047/)

## 7. Ten Simple Rules for Reproducible Research in Jupyter Notebooks

Adam Rule, Amanda Birmingham, Cristal Zuniga, Ilkay Altintas, Shih-Cheng Huang, Rob Knight, Niema Moshiri, Mai H. Nguyen, Sara Brin Rosenthal, Fernando Pérez, Peter W. Rose

*arXiv* (2018-10-13) <https://arxiv.org/abs/1810.08055v1>

## 8. Streamlining Data-Intensive Biology With Workflow Systems

Taylor Reiter, Phillip T. Brooks, Luiz Irber, Shannon E. K. Joslin, Charles M. Reid, Camille Scott, C. Titus Brown, N. Tessa Pierce

*Cold Spring Harbor Laboratory* (2020-11-16) <https://doi.org/gg353v>

DOI: [10.1101/2020.06.30.178673](https://doi.org/10.1101/2020.06.30.178673)

## 9. A Padawan Programmer's Guide to Developing Software Libraries

James T. Yurkovich, Benjamin J. Yurkovich, Andreas Dräger, Bernhard O. Palsson, Zachary A. King

*Cell Systems* (2017-11) <https://doi.org/gg8tqz>

DOI: [10.1016/j.cels.2017.08.003](https://doi.org/10.1016/j.cels.2017.08.003) · PMID: [28988801](https://pubmed.ncbi.nlm.nih.gov/28988801/)



10. **2018 Kaggle Machine Learning & Data Science Survey** <https://kaggle.com/kaggle/kaggle-survey-2018>
11. **Modulecounts** <http://www.modulecounts.com/>
12. **Why scientists are turning to Rust**  
Jeffrey M. Perkel  
*Nature* (2020-12-01) <https://doi.org/ghqc7g>  
DOI: [10.1038/d41586-020-03382-2](https://doi.org/10.1038/d41586-020-03382-2) · PMID: [33262490](https://pubmed.ncbi.nlm.nih.gov/33262490/)
13. **Home - Cookiecutter Data Science** <https://drivendata.github.io/cookiecutter-data-science/>
14. **Reproducible-Science-Curriculum/rr-init**  
Reproducible Science Curriculum  
(2021-02-12) <https://github.com/Reproducible-Science-Curriculum/rr-init>
15. **A Quick Guide to Organizing Computational Biology Projects**  
William Stafford Noble  
*PLoS Computational Biology* (2009-07-31) <https://doi.org/fbbpkn>  
DOI: [10.1371/journal.pcbi.1000424](https://doi.org/10.1371/journal.pcbi.1000424) · PMID: [19649301](https://pubmed.ncbi.nlm.nih.gov/19649301/) · PMCID: [PMC2709440](https://pubmed.ncbi.nlm.nih.gov/PMC2709440/)
16. **Distribution and Reproducibility — Snakemake 5.32.2 documentation**  
<https://snakemake.readthedocs.io/en/stable/snakefiles/deployment.html>
17. **google/styleguide**  
Google  
(2021-02-15) <https://github.com/google/styleguide>
18. **PEP 8 – Style Guide for Python Code**  
Python.org  
<https://www.python.org/dev/peps/pep-0008/>
19. **GitHub: Where the world builds software**  
GitHub  
<https://github.com/>
20. **DevOps Platform Delivered as a Single Application**  
GitLab  
<https://about.gitlab.com/>
21. **Bitbucket | The Git solution for professional teams**  
Atlassian  
*Bitbucket* <https://bitbucket.org/product>
22. **Understanding the GitHub flow · GitHub Guides** <https://guides.github.com/introduction/flow/>
23. **pytest: helps you write better programs — pytest documentation**  
<https://docs.pytest.org/en/stable/>
24. **Codecov - The Leading Code Coverage Solution**  
Codecov  
<https://about.codecov.io/>

25. **Welcome to the tox automation project — tox 3.21.5.dev3 documentation**  
<https://tox.readthedocs.io/en/latest/>
26. **Travis CI - Test and Deploy with Confidence** <https://travis-ci.com/>
27. **Features • GitHub Actions**  
GitHub  
<https://github.com/features/actions>
28. **Reviewing changes in pull requests - GitHub Docs**  
<https://docs.github.com/en/github/collaborating-with-issues-and-pull-requests/reviewing-changes-in-pull-requests>
29. **Crucible Code Review Tool for Git, SVN, Perforce and More**  
Atlassian  
*Atlassian* <https://www.atlassian.com/software/crucible>
30. **Upsource: Code Review and Project Analytics by JetBrains**  
JetBrains  
<https://www.jetbrains.com/upsource/>
31. **The FAIR Guiding Principles for scientific data management and stewardship**  
Mark D. Wilkinson, Michel Dumontier, IJsbrand Jan Aalbersberg, Gabrielle Appleton, Myles Axton, Arie Baak, Niklas Blomberg, Jan-Willem Boiten, Luiz Bonino da Silva Santos, Philip E. Bourne, ... Barend Mons  
*Scientific Data* (2016-03-15) <https://doi.org/bdd4>  
DOI: [10.1038/sdata.2016.18](https://doi.org/10.1038/sdata.2016.18) · PMID: [26978244](https://pubmed.ncbi.nlm.nih.gov/26978244/) · PMCID: [PMC4792175](https://pubmed.ncbi.nlm.nih.gov/PMC4792175/)
32. **Tidy Data**  
Hadley Wickham  
*Journal of Statistical Software* (2014) <https://doi.org/gdm3p7>  
DOI: [10.18637/jss.v059.i10](https://doi.org/10.18637/jss.v059.i10)
33. **Data Version Control • DVC** <https://dvc.org/>
34. **nbviewer** <https://nbviewer.jupyter.org/>
35. **The Binder Project** <https://mybinder.org/>
36. **Google Colaboratory** <https://colab.research.google.com/>
37. **GitHub Diffs & Commenting for Jupyter Notebooks** <https://www.reviewnb.com/>
38. **The WorkflowHub** <https://workflowhub.eu/>
39. **Google Docs: Free Online Documents for Personal Use** <https://www.google.com/docs/about/>
40. **Microsoft 365 with Office apps | Microsoft 365** <https://www.microsoft.com/en-us/microsoft-365>
41. **Open Research Collaboration and Publishing - Authorea** <https://www.authorea.com/>
42. **Overleaf, Online LaTeX Editor** <https://www.overleaf.com>

43. **Manubot - Manuscripts, open and automated** <https://manubot.org>
44. **GitHub CLI**  
GitHub CLI  
<https://cli.github.com/>
45. **About issues - GitHub Docs** <https://docs.github.com/en/github/managing-your-work-on-github/about-issues>
46. **Unit Testing for R** <https://testthat.r-lib.org/>
47. **Flake8: Your Tool For Style Guide Enforcement — flake8 3.8.4 documentation**  
<https://flake8.pycqa.org/en/latest/>
48. **Safety - Security for your Python dependencies** <https://pyup.io/safety/>
49. **Continuous Integration and Delivery**  
CircleCI  
<https://circleci.com/>
50. **How to do a code review**  
eng-practices  
<https://google.github.io/eng-practices/review/reviewer/>
51. **Code Review Guidelines for Humans**  
Philipp Hauer  
*Philipp Hauer's Blog* (2018-07-31) <https://phauer.com/2018/code-review-guidelines/>
52. **Zenodo - Research. Shared.** <https://zenodo.org/>
53. **About**  
figshare  
<https://figshare.com/about>
54. **Git Large File Storage**  
Git Large File Storage  
<https://git-lfs.github.com/>
55. **What is my disk quota? - GitHub Docs** <https://docs.github.com/en/github/managing-large-files/what-is-my-disk-quota>
56. **Open collaborative writing with Manubot**  
Daniel S. Himmelstein, Vincent Rubinetti, David R. Slochower, Dongbo Hu, Venkat S. Malladi, Casey S. Greene, Anthony Gitter  
*PLOS Computational Biology* (2019-06-24) <https://doi.org/c7np>  
DOI: [10.1371/journal.pcbi.1007128](https://doi.org/10.1371/journal.pcbi.1007128) · PMID: [31233491](https://pubmed.ncbi.nlm.nih.gov/31233491/) · PMCID: [PMC6611653](https://pubmed.ncbi.nlm.nih.gov/PMC6611653/)